

# Machine learning from scratch

## Lecture 6: Non-linear models, parameter selection

Alexis Zubiolo

`alexis.zubiolo@gmail.com`

Data Science Team Lead @ Adcash

March 9, 2017

## Course outline

Last time, we reviewed the main ideas behind OLS

# Course outline

Last time, we reviewed the main ideas behind OLS

This lecture will go a bit further by introducing:

- ▶ Non linear models (polynomial kernels)
- ▶ Model evaluation
- ▶ Parameter selection

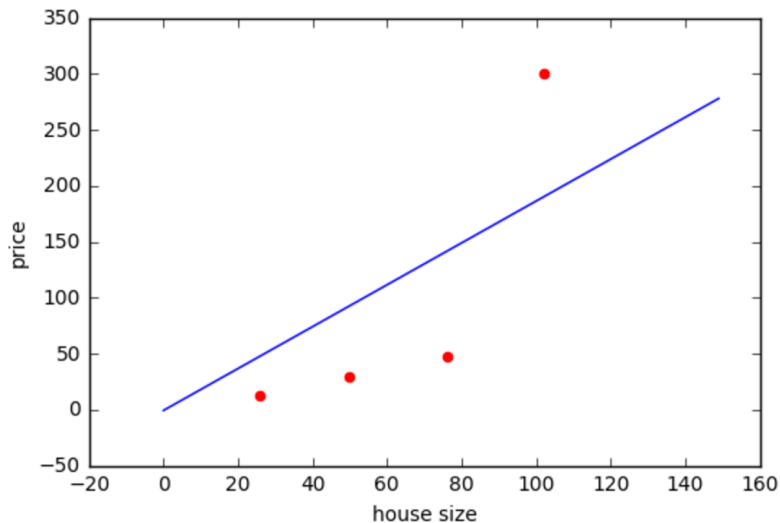
More complex models

## Outliers and overfitting

Recall from previous course: Outliers can ruin the trained model.

## Outliers and overfitting

Recall from previous course: Outliers can ruin the trained model.



## Outliers and overfitting

When this happens, we usually notice that one of the weight is big (in this case,  $\theta_1$ ).

## Outliers and overfitting

When this happens, we usually notice that one of the weight is big (in this case,  $\theta_1$ ).

How to avoid that?



## Outliers and overfitting

When this happens, we usually notice that one of the weight is big (in this case,  $\theta_1$ ).

How to avoid that? By penalizing big weights in the model.

## Outliers and overfitting

When this happens, we usually notice that one of the weight is big (in this case,  $\theta_1$ ).

How to avoid that? By penalizing big weights in the model.

Formally, this consists in rewriting the cost function  $J(\theta)$ . Initially, we had

$$J(\theta) = \frac{1}{2} \sum_{i=1}^n \ell(\hat{y}^{(i)}, y^{(i)})$$

## Outliers and overfitting

When this happens, we usually notice that one of the weight is big (in this case,  $\theta_1$ ).

How to avoid that? By penalizing big weights in the model.

Formally, this consists in rewriting the cost function  $J(\theta)$ . Initially, we had

$$J(\theta) = \frac{1}{2} \sum_{i=1}^n \ell(\hat{y}^{(i)}, y^{(i)})$$

We can add another term  $R(\theta)$  called **regularization**

$$J(\theta) = \frac{1}{2} \sum_{i=1}^n \ell(\hat{y}^{(i)}, y^{(i)}) + \lambda R(\theta)$$

## Outliers and overfitting

When this happens, we usually notice that one of the weight is big (in this case,  $\theta_1$ ).

How to avoid that? By penalizing big weights in the model.

Formally, this consists in rewriting the cost function  $J(\theta)$ . Initially, we had

$$J(\theta) = \frac{1}{2} \sum_{i=1}^n \ell(\hat{y}^{(i)}, y^{(i)})$$

We can add another term  $R(\theta)$  called **regularization**

$$J(\theta) = \frac{1}{2} \sum_{i=1}^n \ell(\hat{y}^{(i)}, y^{(i)}) + \lambda R(\theta)$$

where  $\lambda$  is a **hyper-parameter** that quantifies how much we want to penalize big values of  $\theta$ .

# Outliers and overfitting

In the end,  $J(\theta)$  is as follows:

$$J(\theta) = L(\theta) + \lambda R(\theta)$$

where

- ▶  $L$  is the loss term
- ▶  $R$  is the regularization term

# Outliers and overfitting

In the end,  $J(\theta)$  is as follows:

$$J(\theta) = L(\theta) + \lambda R(\theta)$$

where

- ▶  $L$  is the loss term
- ▶  $R$  is the regularization term

A commonly used regularization term  $R$  is often the squared  $\ell_2$  norm given by

$$R(\theta) = \|\theta\|_2^2 = \sum_{j=1}^d \theta_j^2$$

Model evaluation  
Parameter selection

## Train-test split

As we saw with OLS, ML algorithms usually rely on **many parameters**. How to **tune** them properly given a data set?



# Train-test split

As we saw with OLS, ML algorithms usually rely on **many parameters**. How to **tune** them properly given a data set?

The most commonly used principle is the **train-test split**:

- ▶ **Split the data** into a training set and a test set
- ▶ **Train** on the training set
- ▶ **Test** on the test set

# Train-test split

As we saw with OLS, ML algorithms usually rely on **many parameters**. How to **tune** them properly given a data set?

The most commonly used principle is the **train-test split**:

- ▶ **Split the data** into a training set and a test set
- ▶ **Train** on the training set
- ▶ **Test** on the test set

This is often referred to as **cross-validation**.

# Cross-validation

Standard technique: Hold-out cross-validation:

- ▶ Train on a part of the data (e.g. 70%)
- ▶ Test on the remaining data (e.g. 30%)

# Cross-validation

Standard technique: Hold-out cross-validation:

- ▶ Train on a part of the data (e.g. 70%)
- ▶ Test on the remaining data (e.g. 30%)

Another standard technique:  **$k$ -fold cross-validation**

- ▶ Split the data into  $k$  (equally-sized) folds
- ▶ Remove 1 fold (= test fold)
- ▶ Train on the other folds
- ▶ Test on the removed fold
- ▶ Do it for all the folds

## Small data set

Suppose you have **a small data set**. How to evaluate a classifier in this case?

## Small data set

Suppose you have a **small data set**. How to evaluate a classifier in this case?

$k$ -fold cross-validation? Even with  $k = 2$ , it would make the training set even smaller and make it hard to fit a proper model.

## Small data set

Suppose you have a **small data set**. How to evaluate a classifier in this case?

$k$ -fold cross-validation? Even with  $k = 2$ , it would make the training set even smaller and make it hard to fit a proper model.

Other option: **Leave-one-out** (LOO) cross-validation:

## Small data set

Suppose you have a **small data set**. How to evaluate a classifier in this case?

$k$ -fold cross-validation? Even with  $k = 2$ , it would make the training set even smaller and make it hard to fit a proper model.

Other option: **Leave-one-out** (LOO) cross-validation:

- ▶ **Remove 1 sample** from the data set
- ▶ Train on **all the other samples**
- ▶ Test on the sample you've removed
- ▶ **Evaluate** the prediction
- ▶ Do it **for each sample of the data set**
- ▶ **Aggregate** the evaluations



## Small data set

Suppose you have a **small data set**. How to evaluate a classifier in this case?

$k$ -fold cross-validation? Even with  $k = 2$ , it would make the training set even smaller and make it hard to fit a proper model.

Other option: **Leave-one-out** (LOO) cross-validation:

- ▶ **Remove 1 sample** from the data set
- ▶ Train on **all the other samples**
- ▶ Test on the sample you've removed
- ▶ **Evaluate** the prediction
- ▶ Do it **for each sample of the data set**
- ▶ **Aggregate** the evaluations

**Remark:** This could lead to many iterations even if the data set is small.

## Small data set

Suppose you have a **small data set**. How to evaluate a classifier in this case?

$k$ -fold cross-validation? Even with  $k = 2$ , it would make the training set even smaller and make it hard to fit a proper model.

Other option: **Leave-one-out** (LOO) cross-validation:

- ▶ **Remove 1 sample** from the data set
- ▶ Train on **all the other samples**
- ▶ Test on the sample you've removed
- ▶ **Evaluate** the prediction
- ▶ Do it **for each sample of the data set**
- ▶ **Aggregate** the evaluations

**Remark:** This could lead to many iterations even if the data set is small.

**Alternative:** Leave- $p$ -out (LPO). LOO is LPO with  $p = 1$ .

## Parameter selection

One of the goals of model evaluation is to select a **good model**.

## Parameter selection

One of the goals of model evaluation is to select a **good model**.

Hence we want to choose one (or several) parameters. How do we do?

# Parameter selection

One of the goals of model evaluation is to select a **good model**.

Hence we want to choose one (or several) parameters. How do we do?

Most classic way: a **grid-search**

- ▶ For each parameter, define a set of possible values
- ▶ For each parameter combination, train/test the model
- ▶ Pick the parameter combination which gives best results

# Parameter selection

One of the goals of model evaluation is to select a **good model**.

Hence we want to choose one (or several) parameters. How do we do?

Most classic way: a **grid-search**

- ▶ For each parameter, define a set of possible values
- ▶ For each parameter combination, train/test the model
- ▶ Pick the parameter combination which gives best results

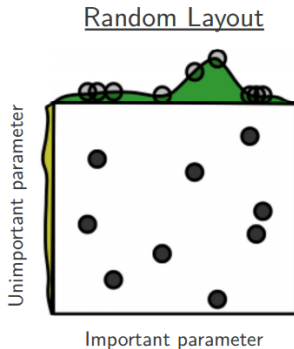
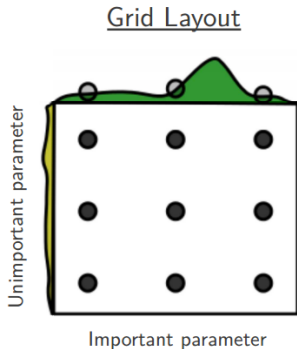
**Important note:** Hyper-parameter ranges vary a lot from an application to another. It is **data-dependent**.

## Grid search vs random search

Random search is a more and more popular alternative to grid search, especially in this case:

## Grid search vs random search

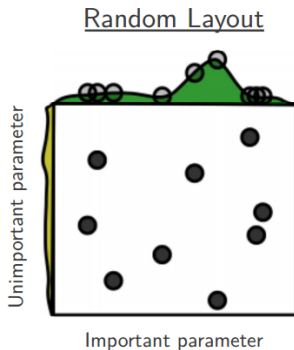
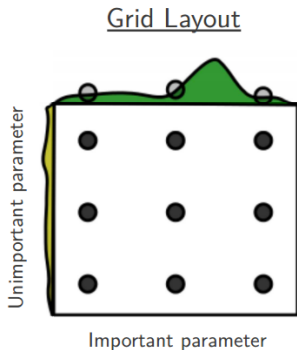
Random search is a more and more popular alternative to grid search, especially in this case:





## Grid search vs random search

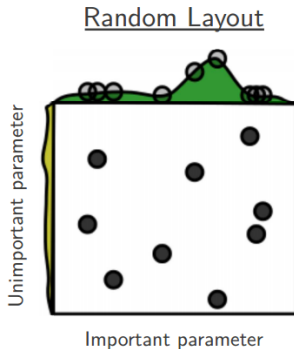
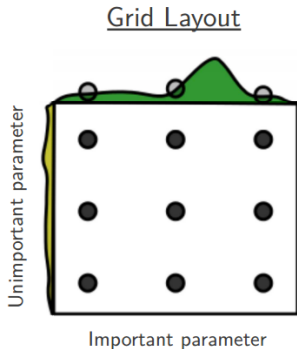
Random search is a more and more popular alternative to grid search, especially in this case:



In any case, you need to guess upper/lower bounds on the parameters.

## Grid search vs random search

Random search is a more and more popular alternative to grid search, especially in this case:



In any case, you need to guess upper/lower bounds on the parameters.

**Practical note:** Each parameter combination can be trained/tested separately => possibility to distribute the tasks

# Conclusion

In this lecture, we've seen:

- ▶ How to avoid overfitting by regularizing the model
- ▶ How to evaluate models
- ▶ How to tune parameters automatically

# Conclusion

In this lecture, we've seen:

- ▶ How to avoid overfitting by regularizing the model
- ▶ How to evaluate models
- ▶ How to tune parameters automatically

During the next lecture, we will work on implementing regularization to the OLS algorithm and cross-validating it and switch to classification if the time allows it.

Thank you! Questions?