

# Machine learning from scratch

## Lecture 7: Classification

Alexis Zubiolo

`alexis.zubiolo@gmail.com`

Data Science Team Lead @ Adcash

March 16, 2017

# Introduction

**Classification** and **regression** are the 2 main aspects of supervised machine learning.

# Introduction

**Classification** and **regression** are the 2 main aspects of supervised machine learning.

**Reminder:** In classification,  $y$  takes **discrete values: 1 value per class**. Typically, we note  $y \in \{1, 2, \dots, k\}$  where  $k$  is the number of classes.

# Introduction

**Classification** and **regression** are the 2 main aspects of supervised machine learning.

**Reminder:** In classification,  $y$  takes **discrete values: 1 value per class**. Typically, we note  $y \in \{1, 2, \dots, k\}$  where  $k$  is the number of classes.

**Example:**

- ▶ For digit recognition, there are **10 classes** (one per digit: 0, 1, 2, ..., 9)
- ▶ For letter recognition, the number of classes depends on the alphabet (e.g. 26 for the latin alphabet).

# Support Vector Machines

**Support Vector Machines** (SVM) is one of the most popular classification algorithms. What we will see in this course is just one way to approach it. Actually, it can be solved/approached in many (equivalent) ways.

# Support Vector Machines

**Support Vector Machines** (SVM) is one of the most popular classification algorithms. What we will see in this course is just one way to approach it. Actually, it can be solved/approached in many (equivalent) ways.

As usual, we will define a dataset  $\{(\mathbf{x}^{(i)}, y^{(i)}), i = 1, \dots, n\}$ . Here,  $y^{(i)} \in \{1, \dots, k\}$  is the class.

# Support Vector Machines

**Support Vector Machines** (SVM) is one of the most popular classification algorithms. What we will see in this course is just one way to approach it. Actually, it can be solved/approached in many (equivalent) ways.

As usual, we will define a dataset  $\{(\mathbf{x}^{(i)}, y^{(i)}), i = 1, \dots, n\}$ . Here,  $y^{(i)} \in \{1, \dots, k\}$  is the class.

One way to formulate the SVM model is by defining **one linear models per class**:  $\theta_1, \theta_2, \dots, \theta_k$ . The goal is to have  $\theta_j^T \mathbf{x}$  measure the confidence of  $\mathbf{x}$  having  $j$  as a label. This way, the predicted class of  $\mathbf{x}$  will be the  $j$  that maximizes  $\theta^T \mathbf{x}$ .

# Support Vector Machines

Let us consider a single sample  $(\mathbf{x}, y)$ . The SVM algorithm will define weights  $\theta_1, \theta_2, \dots, \theta_k$  for each class.



# Support Vector Machines

Let us consider a single sample  $(\mathbf{x}, y)$ . The SVM algorithm will define weights  $\theta_1, \theta_2, \dots, \theta_k$  for each class.

Hence, we can define scores  $s_j = \theta_j^T \mathbf{x}$  for each class  $j$ . If the SVM algorithm makes no mistake on the sample  $(\mathbf{x}, y)$ ,  $s_y$  **should be the highest value**.

# Hinge loss

SVM relies on the **hinge loss**, that is

$$\ell(s_j, s_y) = \max(0, s_j - s_y + \Delta)$$

where:

- ▶  $s_j$  is the score for some class  $j \neq y$
- ▶  $s_y$  is the score for the true class  $y$
- ▶  $\Delta$  is a **hyper-parameter** that quantifies by how much we want  $s_y$  to be bigger than  $s_j$  for  $j \neq y$

## Hinge loss

SVM relies on the **hinge loss**, that is

$$\ell(s_j, s_y) = \max(0, s_j - s_y + \Delta)$$

where:

- ▶  $s_j$  is the score for some class  $j \neq y$
- ▶  $s_y$  is the score for the true class  $y$
- ▶  $\Delta$  is a **hyper-parameter** that quantifies by how much we want  $s_y$  to be bigger than  $s_j$  for  $j \neq y$

This loss has to be computed for all the classes **other than the actual class**  $y$ . In the end, the total loss for the sample  $(\mathbf{x}, y)$  is the sum of the loss for all the  $j \neq y$ :

$$L = \sum_{j \neq y} \ell(s_j, s_y) = \sum_{j \neq y} \max(0, s_j - s_y + \Delta)$$

## Illustration on a simple example

**Example:** Suppose we have 3 classes (1, 2 and 3). We get a training sample  $(\mathbf{x}, y)$  within the first class (that is,  $y = 0$ ). We have an SVM algorithm with hyperparameter  $\Delta = 10$  that returns scores of 13 for class 0,  $-7$  for class 1, and 11 for class 2.

**Question:** What is the total loss for this sample?

## Illustration on a simple example

**Example:** Suppose we have 3 classes (1, 2 and 3). We get a training sample  $(\mathbf{x}, y)$  within the first class (that is,  $y = 0$ ). We have an SVM algorithm with hyperparameter  $\Delta = 10$  that returns scores of 13 for class 0,  $-7$  for class 1, and 11 for class 2.

**Question:** What is the total loss for this sample?

**Solution:**

$$\begin{aligned} L &= \ell(s_1, s_0) + \ell(s_2, s_0) \\ &= \max(0, -7 - 13 + 10) + \max(0, 11 - 13 + 10) \\ &= \max(0, -10) + \max(0, +8) \\ &= 0 + 8 \\ &= 8 \end{aligned}$$

## Optimization problem formulation

We saw that for a single sample  $(\mathbf{x}, y)$ , the training loss is:

$$L = \sum_{j \neq y} \ell(s_j, s_y) = \sum_{j \neq y} \max(0, s_j - s_y + \Delta)$$

where  $s_j = \theta_j^T \mathbf{x}$  and  $s_y = \theta_y^T \mathbf{x}$ .

## Optimization problem formulation

We saw that for a single sample  $(\mathbf{x}, y)$ , the training loss is:

$$L = \sum_{j \neq y} \ell(s_j, s_y) = \sum_{j \neq y} \max(0, s_j - s_y + \Delta)$$

where  $s_j = \theta_j^T \mathbf{x}$  and  $s_y = \theta_y^T \mathbf{x}$ .

As for the linear regression, when given a training set  $\{(\mathbf{x}^{(i)}, y^{(i)}), i = 1, \dots, n\}$ , we just need to sum the losses for all the training samples:

$$\begin{aligned} J(\theta_1, \theta_2, \dots, \theta_k) &= \sum_{i=1}^n L_i \\ &= \sum_{i=1}^n \sum_{j \neq y_i} \max(0, \theta_j^T \mathbf{x}^{(i)} - \theta_{y_i}^T \mathbf{x}^{(i)} + \Delta) \end{aligned}$$

## Overfitting and regularization

As for regression tasks, **outliers** can have a dramatic impact on the model we train. We saw that it usually resulted in high weights and that **adding a regularization term** could help.

$$R(\theta_1, \theta_2, \dots, \theta_k) = \sum_{j=1}^k \|\theta_j\|_2^2$$



## Overfitting and regularization

As for regression tasks, **outliers** can have a dramatic impact on the model we train. We saw that it usually resulted in high weights and that **adding a regularization term** could help.

$$R(\theta_1, \theta_2, \dots, \theta_k) = \sum_{j=1}^k \|\theta_j\|_2^2$$

We can do the same for the SVM problem:

$$J(\theta_1, \theta_2, \dots, \theta_k) = \sum_{i=1}^n \sum_{j \neq y_i} \max(0, \theta_j^T \mathbf{x}^{(i)} - \theta_{y_i}^T \mathbf{x}^{(i)} + \Delta) + \lambda \sum_{j=1}^k \|\theta_j\|_2^2$$

where

$$\|\theta_j\|_2^2 = \sum_{m=1}^d \theta_{j,m}^2$$

( $\theta_j$  is a vector, and  $\theta_{j,m}$  its  $m$ th value.)

# Optimization

We have defined a cost function  $J$  we would like to minimize. Again, we can apply an algorithm like gradient descent to find its optimal value. For this, we need to compute  $J$ 's **gradients**.

# Optimization

We have defined a cost function  $J$  we would like to minimize. Again, we can apply an algorithm like gradient descent to find its optimal value. For this, we need to compute  $J$ 's **gradients**.

In this case, there will be **several gradients** (1 per class because we have 1  $\theta_j$  vector per class). We will note  $\nabla_{\theta_j} J$  the gradient with respect to  $\theta_j$ .

# Optimization

We have defined a cost function  $J$  we would like to minimize. Again, we can apply an algorithm like gradient descent to find its optimal value. For this, we need to compute  $J$ 's **gradients**.

In this case, there will be **several gradients** (1 per class because we have 1  $\theta_j$  vector per class). We will note  $\nabla_{\theta_j} J$  the gradient with respect to  $\theta_j$ .

Due to the gradient linearity, we can compute the gradient of each component **separately**. Hence, we need to compute:

- ▶  $L$ 's gradients
- ▶  $R$ 's gradients

## $L$ 's gradients

**Reminder:** For a sample  $(\mathbf{x}, y)$ ,  $L$  is given by:

$$L = \sum_{j \neq y} \max(0, \theta_j^T \mathbf{x} - \theta_y^T \mathbf{x} + \Delta)$$

## $L$ 's gradients

**Reminder:** For a sample  $(\mathbf{x}, y)$ ,  $L$  is given by:

$$L = \sum_{j \neq y} \max(0, \theta_j^T \mathbf{x} - \theta_y^T \mathbf{x} + \Delta)$$

There are 2 different cases:

## $L$ 's gradients

**Reminder:** For a sample  $(\mathbf{x}, y)$ ,  $L$  is given by:

$$L = \sum_{j \neq y} \max(0, \theta_j^T \mathbf{x} - \theta_y^T \mathbf{x} + \Delta)$$

There are 2 different cases:

►  $j \neq y$ :

$$\nabla_{\theta_j} L = \begin{cases} \mathbf{x} & \text{if } \theta_j^T \mathbf{x} - \theta_y^T \mathbf{x} + \Delta > 0 \\ 0 & \text{otherwise.} \end{cases}$$

## $L$ 's gradients

**Reminder:** For a sample  $(\mathbf{x}, y)$ ,  $L$  is given by:

$$L = \sum_{j \neq y} \max(0, \theta_j^T \mathbf{x} - \theta_y^T \mathbf{x} + \Delta)$$

There are 2 different cases:

►  $j \neq y$ :

$$\nabla_{\theta_j} L = \begin{cases} \mathbf{x} & \text{if } \theta_j^T \mathbf{x} - \theta_y^T \mathbf{x} + \Delta > 0 \\ 0 & \text{otherwise.} \end{cases}$$

►  $j = y$ :

$$\nabla_{\theta_j} L = p \mathbf{x}$$

where  $p$  is the **number of times the desired margin is not satisfied**, that is the number of  $j \neq y$  such that

$$\theta_j^T \mathbf{x} - \theta_y^T \mathbf{x} + \Delta > 0$$



## $R$ 's gradients

**Reminder:**  $R$  is given by:

$$R(\theta_1, \theta_2, \dots, \theta_k) = \sum_{j=1}^k \|\theta_j\|_2^2$$

**Question:** What is  $\nabla_{\theta_j} R$ ?

## $R$ 's gradients

**Reminder:**  $R$  is given by:

$$R(\theta_1, \theta_2, \dots, \theta_k) = \sum_{j=1}^k \|\theta_j\|_2^2$$

**Question:** What is  $\nabla_{\theta_j} R$ ?

**Answer:**

$$\nabla_{\theta_j} R = 2\theta_j$$

## $R$ 's gradients

**Reminder:**  $R$  is given by:

$$R(\theta_1, \theta_2, \dots, \theta_k) = \sum_{j=1}^k \|\theta_j\|_2^2$$

**Question:** What is  $\nabla_{\theta_j} R$ ?

**Answer:**

$$\nabla_{\theta_j} R = 2\theta_j$$

**Conclusion:** We have the analytic expression of the gradients of  $L$  and  $R$  and are able to apply the gradient descent (stochastic or batch) to it.

# Conclusion

In this lecture, we've seen Classification models with the example of SVMs:

- ▶ How to define the problem
- ▶ How to compute the gradient
- ▶ How to train an SVM model for multiclass classification

# Conclusion

In this lecture, we've seen Classification models with the example of SVMs:

- ▶ How to define the problem
- ▶ How to compute the gradient
- ▶ How to train an SVM model for multiclass classification

Next time, we'll go a bit further and have a **practical session** about classification.

Thank you! Questions?